

- In linear regression, assuming Gaussian noise leads to  $L_2$  loss, because maximizing Gaussian likelihood is equivalent to minimizing squared error.
- Ridge regression corresponds to MAP estimation with Gaussian prior on the weights.
- Cross validation is commonly used to choose hyperparameters and estimate test error
- in ridge, as  $\lambda \uparrow$ , typically bias  $\uparrow$  and variance  $\downarrow$
- $L_0$ : stepwise search (counts features),  $L_1$ : sparse solutions,  $L_2$ : Shrinks weights but none to 0.
- the softmax function maps a score vector to a probability distribution
- A standard way of picking RBF centers is to run k-means and use the cluster centers
- Gaussian RBFs are typically local. polynomial basis functions are global?
- If an event is certain, the entropy  $H(X) = 0$
- In NNs, max pooling reduces computation and adds some translation invariance
- PCA minimizes the reconstruction error!
- the distributional similarity hypothesis used for word embeddings states: "words occurring in similar contexts tend to have similar meanings"
- Autoencoders learn to reconstruct inputs through an encoder-decoder.
- A linear autoencoder approximates PCA
- A denoising autoencoder is trained to reconstruct a clean signal from a corrupted input
- k means minimizes the within-cluster squared reconstruction error
- Naive Bayes assumes that words are indep given the topic / class
- If you don't observe the topic in Naive Bayes, you can use MLE instead
- Two sets of nodes  $X$  and  $Y$  are d-separated by a set of nodes  $Z$  if  $X$  and  $Y$  are conditionally independent given  $Z$ :  $X \perp Y | Z$
- $L_0 \nsubseteq L_{1/2}$  are not convex,  $L_1 \subseteq L_2$  are convex.
- kmeans assumes all clusters are roughly the same size (inductive bias)
- CNN inductive biases: locality, translational invariance (
- RL: MDPs generalize Markov Models. POMDPs generalize HMMs. "MDPs generalize to NNets"

"In the limit, ridge reg shrinks the params to 0  
hence 0 variance" OLS linear regression  
is unbiased in  $\hat{\theta}$  or  $\hat{y}$ , ridge makes it biased

How to solve

$$p=2 : (X^T X + \lambda I)^{-1} X^T Y$$

$p=1$ : Gradient descent

$p=0$  search (stepwise or streamwise)



$$L_2 \text{ regression} = \underset{w}{\operatorname{argmin}} \|y - Xw\|_2 = \text{OLS?}$$

$L_2$  most heavily shrinks large weights

- $L_0$  penalty most strongly encourages weights to be 0
- $L_0$  penalty is scale invariant,  $L_1 \neq L_2$  are not

Elastic net, cross, = ridge + lasso

kernel regression v/s k-nn:




If we have case B, we rather use k-nn

If we have case A, we rather use kernel regression!


relationship k means knn.

HL: i) components: Rep, Loss, Opt ; ii) types: SL, UL, RL ;  
 iii) models: generative p(m) } parametric  
 discriminative p(y|m) } non-parametric knn, decision trees, kernel regression  
 semi-parametric



k-MV can do both regression & classification:  $k \uparrow \Rightarrow$  model complexity  $\downarrow$   
 non-parametric but has two hyperparameters  $k$  and  $d(\cdot, \cdot)$   
 algo: 1) find  $k$  nearest neighbors according to  $d(\cdot, \cdot)$  distance metric  
 2)  $\hat{y}(m)$  = majority label or average  $y$ 's of those  $k$  neighbors  
 distance properties: i)  $d(x, y) \geq 0$ ; ii)  $d(x, y) = 0 \Leftrightarrow x = y$ ; iii)  $d(x, y) = d(y, x)$ ; iv)  $d(x, z) \leq d(x, y) + d(y, z)$

$L_p$  norm:  $\|x\|_p = \left(\sum_{i=1}^n |x_i|^p\right)^{1/p}$   $\rightarrow$  props: i)  $\|x\| = 0 \Leftrightarrow x$  is the 0 vector ii)  $\|a x\|_p = |a| \|x\|_p$  (homogeneity)  
 iii)  $\|x+y\|_p \leq \|x\|_p + \|y\|_p$  (triangle inequality or subadditivity) iv)  $p \geq 1$



$L_0$  pseudo-norm =  $\|x\|_0$  = the number of non-zero components of  $x$   
 $L_\infty = \|x\|_\infty = \max |x_i|$  = largest absolute component of  $x$   
 curse of dimensionality: as dimension  $\uparrow$ , variance of distances shrink relative to their mean  $\Rightarrow$  points are equally close

2a) GO:  $\theta_{i+1} = \theta_i - \eta \nabla_{\theta} \mathcal{L}(\theta)$ ; mini-batch: "update model every  $k$  observations"; SGO: "mini batch with  $k=1$ " ( $k$ =batch size)  
 momentum: "continue moving in same direction"; RMS Prop: "small (large) steps for large (small) gradients"; Adam: "momentum + RMS Prop"  
 Adagrad: uses a per-feature learning rate; Coordinate descent: "use when in  $L_1$  world"

3a) Likelihood prior  $P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)}$ ; HLE:  $\arg \max_{\theta} P(D|\theta)$   $\rightarrow$  solve  $\rightarrow \log(\cdot) = \frac{\partial}{\partial \theta} = 0$  NLE & MAP both consistent; NLE overfits more than MAP  
 posterior marginal  $P(\theta)$ ; MAP:  $\arg \max_{\theta} P(D|\theta)P(\theta)$  NLE = MAP with uniform prior;  $X^T X$  singular  $\Rightarrow$  no NLE  
 $\hat{w}_{HLE} = \frac{\sum x_i y_i}{\sum x_i^2}$ ;  $\hat{w}_{HLE} = (X^T X)^{-1} X^T Y$ ;  $\hat{w}_{MAP} = (X^T X + I)^{-1} X^T Y$   $\hat{w}_{OLS} = \hat{w}_{HLE}$  w/ Gaussian noise;  $\hat{w}_{ridge} = \hat{w}_{MAP}$  w/ Gaussian prior on  $w$   
 simple model:  $b \uparrow v \downarrow$ ; complex model:  $b \downarrow v \uparrow$

3b) The simpler (more complex) the model, the less (more) well it's gonna fit the data, but the less (more) it's gonna vary for  $\neq$  datasets (i.e. the more (less) stable it's gonna be)  
 Test error = Train error + Model complexity = Variance + Bias + Noise  
 K-Fold CV: used to pick hyperparams & estimate test error; if  $K=n \Rightarrow$  LOOCV  
 bias  $(\hat{\theta}) = E[\hat{\theta} - \theta]$ ; var  $(\hat{\theta}) = E[(\hat{\theta} - E[\hat{\theta}])^2]$   
 adding penalties to OLS:  $\|y - Xw\|_2^2 + \lambda \|w\|_p^p$   
 $\Leftrightarrow$  MAP with Gaussian ( $p=2$ ), Laplace ( $p=1$ ) priors  
 $p=2$  (ridge) convex, makes  $w$ 's smaller but never 0  
 $p=1$  LASSO convex, drive some  $w$ 's to 0 (feature selection)  
 $p=0$  (stepwise) nonconvex, requires search, but stronger than  $L_2$




stepwise regression: look at all remaining features at each step and pick the best to add (n-greedy)  
 streamwise regression: look at one feature at a time, in a fixed order, immediately accept or reject, no revisits (more greedy)

4a) Logistic regression for binary classification, uses the Sigmoid function:  $\sigma(z) = \frac{1}{1+e^{-z}}$ ; log odds:  $\frac{P(y=1|x,w)}{P(y=0|x,w)} = w^T x$  if  $w=0 \Rightarrow$  hyperplane decision bound  
 Multi-class classification for  $K$  classes uses  $p(y=k|x, \theta_1, \dots, \theta_K) = \frac{e^{\theta_k^T x}}{\sum_{k=1}^K e^{\theta_k^T x}}$  (softmax function)  $\rightarrow$  maps a vector to a probability distro  
 4 predict class as most probable label "kernel width"  
 RBF:  $k\left(\frac{\|x - x_c\|_2}{c}\right) = \exp\left\{-\frac{\|x - x_c\|_2^2}{c^2}\right\}$   $c$  small  $\Rightarrow$  overfit,  $c$  big  $\Rightarrow$  underfit  
 4 to find kernel centers, do  $k$ -means and take cluster centers  
 i)  $d < p$  basis vectors  $\Rightarrow$  dimensionality reduction  
 ii)  $d > p \Rightarrow$  dimensionality increase (make non-linear prob linear)  
 iii)  $d = n \Rightarrow$  switch to dual representation

4b) a kernel function measures the similarity between two points and induces an SPSD matrix  $K_{ij} := k(x_i, x_j)$  kernel:  $\hat{y}(x) = \frac{\sum K(x, x_i) y_i}{\sum K(x, x_i)}$   $\rightarrow$  non parametric  
 "kernels assume similarity is the same everywhere"  
 SVM: used for classification with hinge loss =  $\max(0, 1 - y * f(m))$  "Points on correct side of margin don't contribute to hinge loss, so loss entirely in terms of points inside margin or on wrong side of it"  
 kernel class:  $\hat{y}(x) = \text{sign}(\sum K(x, x_i) y_i)$

5a) entropy of a rv  $X$ ;  $H(X) \in [0, \infty)$  but  $\in [0, \infty]$  if  $X$  binary  
 $H(X) = -\sum_{i=1}^n p_i \log_2(p_i)$ ;  $m = \#$  events  $E_i$  and  $p_i = P(E_i)$   
 = "average conditional entropy of  $x$ "  
 conditional entropy:  $H(Y|X) = \sum_{i=1}^n P(X=E_i) H(Y|X=E_i)$ ; Information gain:  $IG(Y|X) = H(Y) - H(Y|X)$ ; cross entropy:  $H(P, Q) = -\sum p(x) \log q(x)$   
 KL divergence:  $D_{KL}(P||Q) = \sum p(x_i) \log \frac{p(x_i)}{q(x_i)} = \dots = H(P, Q) - H(P) \geq 0$ ; non-symmetric, measures difference, and  $D_{KL}(P||Q) = 0 \Leftrightarrow P=Q$   
 decision tree algo: Recursive partition to maximize IG. if overfit  $\rightarrow$  pruning, max depth, or other



- two events are equally likely  $\Rightarrow H(X) = 1$
- $H(X) = 0 \Rightarrow$  an event is certain
- better Model  $\Rightarrow$  lower  $H(Y|X)$



**Ensembles:** combine many weak models into a strong one. **bagging:** draw bootstrap sample with replacement, train weak model. **Random Forests:** Bagging + random feature selection (no overfit, no underfit) → repeat for T models. **stagewise regression:** sequentially learn the weights  $\alpha_t$  assigned to models, never readjust; **boosting:** "give more weight to points we got wrong". **Adaboost:** minimizes exponential loss, learns exponentially fast, doesn't overfit; **Gradient(T) Boosting:** generalization of ada for any loss. **regularization for GB:**  $\eta$  (shrinkage); tree depth; # boosting stages, ... reduces variance but also bias unlike RF

**8a) SVD:**  $X = UDV^T$ ; cols of U and V are left and right singular vectors and diag elems of D (diag) are the singular values of X.  $\|A\|_F = \sqrt{\sum_i |a_{ij}|^2} = \sqrt{\text{trace}(A^T A)} = \sqrt{\sum_i \sigma_i^2(A)}$  not  $V^T$ !  $X = ZV^T$  where  $Z = UD$  = scores &  $V$  = loadings, PC's, PC direction/components. **PCA:** minimizes distortion  $\min_{Z,V} \|X - ZV^T\|_F^2 \leq 7$  maximizes variance  $\min_{Z,V} \|X - ZV^T\|_F^2 + \lambda_1 \|Z\|_1 + \lambda_2 \|V\|_1 \rightarrow$  convex in Z|V or V|Z. **Sparse PCA:** introduce  $L_1$  penalty on scores or loadings  $\Rightarrow$  no longer eigenvalues, must solve using alternating GD. **Moore-Penrose pseudo inverse of X:**  $X^+ = VD^{-1}U^T$ , thin SVD:  $X^+ \sim V_k D_k^{-1} U_k^T$  (cheap to k singular vals/vectors) good approx of inverse! **Power Method:** finds PC1 by repeatedly multiplying a random vector by  $X^T X$ , causing it to align with the dominant eigenvector of  $X^T X$ : PC1.

**8b) PCR:** using PCA scores as regression features → 1) do PCA on X get loadings & scores; 2) do OLS using scores:  $w = (Z^T Z)^{-1} Z^T y$ . 3) compute score  $z = V^T x$  and predict  $\hat{y} = w^T z = w^T V^T x$ . PCA removes collinearity and noise  $\Rightarrow$  regression more stable. **semi-supervised PCR:** use lots of unlabeled data to find good PCA representation, then regress on the labeled data. **PCA for visualization:** PCA reduces data to the top 2 PCs to plot and visually inspect patterns, clusters and relationships. each PC has an eigenval  $\lambda_i$  and % variance explained by  $PC_i = \frac{\lambda_i}{\sum \lambda_j}$ .  $U$  = word embeddings (context-free);  $V$  = context embeddings (context sensitive).

**8c) autoencoders:** NNets taking X as input & output and force info through bottleneck/sparsity to learn an embedding. **denoising AE:** Corrupt input X with noise, reconstruct X cleanly  $\Rightarrow$  learn robust features. Also, **stacked AE**, useful for SL. **ICA:** Given X, find W s.t.  $X = SW$  where  $s_j \perp$ , not just uncorrelated. ICA doesn't have simple SVD solution.  $\rightarrow$  reconstruct  $X \sim SW^T$  ( $S \sim$  PCA scores,  $w^+ \sim$  loadings).  $KL(y_1, \dots, y_m) = KL(p(y_1, \dots, y_m) \| p(y_1) \dots p(y_m))$ .  $\rightarrow$  ICA aims for  $\perp$  not variance! It does disentanglement, as well as reconstruction.

**9a) Confusion matrix:**

|             |            |        |
|-------------|------------|--------|
|             | Prediction |        |
|             | P          | N      |
| Observation | P          | TP, FN |
|             | N          | FP, TN |

**Accuracy** =  $\frac{TP+TN}{TP+TN+FP+FN}$  misleading under imbalance; **Precision** =  $\frac{TP}{TP+FP} = p(\text{yes}|\text{predicted yes})$  measures correctness of positive predictions; **Specificity** =  $\frac{TN}{TN+FP} = p(\text{predicted no}|\text{no})$  measures correctness on negatives; **Recall = Sensitivity** =  $\frac{TP}{TP+FN} = p(\text{predicted yes}|\text{yes})$  measures catching positives; **1 - False positive rate** =  $\frac{TN}{TN+FP}$  min proba to call something positive; **F1 score** =  $\frac{2PR}{P+R}$  if both P↑ R↑  $\Rightarrow$  F1↑

**Precision/Recall trade-off:** threshold  $\uparrow \Rightarrow$  predicted positives  $\downarrow$ , P↑, R↓; threshold  $\downarrow \Rightarrow$  predicted positives  $\uparrow$ , R↑, P↓.  $\Rightarrow$  P↑  $\Rightarrow$  R↓ (vice versa), only way to increase both is via a better model, not a threshold. For medical diagnosis we want Recall↑. For spam email detection we want specificity↑.

**ROC:** examines classifier performance across all thresholds.  $\rightarrow$  evaluates ranking ability, not final classification. **AUC** =  $\int_0^1 \text{ROC curve}$   $\Rightarrow$  perfect 1, 0.5  $\Rightarrow$  random.

**10a) K-means:** 1) pick K cluster centroids at random; 2) assign points to nearest centroid; 3) set centroid to mean of examples assigned to it; 4) repeat from step 2) until convergence.  $\rightarrow$  usually by dist.  $\rightarrow$  alternates between assigning points & updating centroids. K-means minimizes reconstruction error:  $J(\mu, \Gamma) = \sum_{i=1}^n \sum_{k=1}^K \mathbb{1}_{i \in \Gamma_k} \|x_i - \mu_k\|_2^2$ .

$\mu_k = \frac{\sum_{i \in \Gamma_k} x_i}{|\Gamma_k|}$   $\Rightarrow$  alternating minimization, guaranteed to decrease loss at each step, but not guaranteed to find global optimum (initialisation matters!) Hence, run k-means many times to avoid poor solutions. **K-means++:** picks diverse starting points by weighting points by squared distance for chosen centers. **k-means assumes:** i) hard clustering (1 cluster/point) ii) clusters are spherical & equal size iii) cluster proba unif:  $\pi_k = \frac{1}{K}$ . **k-means is a special case of GMMs.** **GMMs** assumes each cluster is generated from a Gaussian distro with params  $(\mu_k, \Sigma_k, \pi_k)$ . **GMM is a soft clustering algo.** **k-means = GMM with equal mixing weights  $\pi_k = \frac{1}{K}$  and  $\Sigma_k = \sigma^2 I$  (shared spherical covariance).** **EM algo** alternates 1) estimate soft cluster membership  $p(z_i = k | x_i)$ ; 2) update parameters (EM only finds local max).

**11a) GTH:** to generate each point: 1) pick a cluster  $z$  w.p.  $\pi$ ; 2)  $p(x) = \sum_k \pi_k \mathcal{N}(x | \mu_k, \Sigma_k)$  where  $x \sim \mathcal{N}(\mu_z, \Sigma_z)$ . **GTH uses EM** because cluster assignment are latent:  $E$ :  $p(z=k|x)$ ,  $M$ : update  $\mu_k, \Sigma_k, \pi_k$ . **NB:** assumes words are conditionally  $\perp$  given topic ( $\rightarrow$  "naive"). If topic labels are known  $\rightarrow$  count, else  $\rightarrow$  EM.  $\rightarrow$  infer topic of document;  $\rightarrow$  Estimate word distribution. **Limits:** i) 1 topic/doc ii) no word dependencies; iii) unrealistic linguistically, hence  $\Rightarrow$  **LDA:** for each doc, sample topic proportion  $\theta_d \sim \text{Dirichlet}(\alpha)$ ; for each word choose topic  $z_n \sim \text{multinomial}(\theta_d)$ ; choose word  $w_n \sim \text{multinomial}(\beta_{z_n})$  ( $\neq$  w/NS: Each word can have its own topic). **HMMs:** 1) sample init state  $s_0 \sim \pi$  2) repeat: i) Emit observation  $x_t \sim p(x_t | s_t)$  ii) transition  $s_{t+1} \sim p(s_{t+1} | s_t)$

HMMs are **dynamic generative models**. For estimation **E**: estimate hidden state; **I**: estimate emission and transition probabilities

**Belief Nets**: DAG, encodes conditional  $\perp$ , factorizes joint:  $p(x_1, \dots, x_n) = \prod p(x_i | \text{Parents}(x_i))$  (key:  $\perp$  assumptions) graph structure

**LMA**: A node is  $\perp$  of all non-descendants given its parents:  $X \perp \text{nondec}_X | \text{Parents}_X$

props: (ii)  $X \perp Y, W | Z \Rightarrow X \perp Y | Z$ ; (iii)  $X \perp Y, W | Z \Rightarrow X \perp Y | Z, W$  (iv)  $(X \perp W | Z), (X \perp Y | Z) \Rightarrow X \perp Y, W | Z$

i)  $X \perp Y | Z \Rightarrow Y \perp X | Z$  (blue parents  $\subset$  non-descendants) •  $X \not\perp Y$  are **d-separated** by  $Z$  if every path between  $X$  and  $Y$  is blocked given  $Z$  (if  $\nexists$  active trail between  $X$  and  $Y$ )  $\Rightarrow X \perp Y | Z$

**Belief nets provide a compact representation of joint probability distributions**

- Guaranteed to be a consistent specification
- Graphically show conditional independence
- Can be extended to include "action" nodes

**Networks that capture causality tend to be sparser**

**Estimating belief nets is easy if everything is observable**

- But requires serious search if the network structure is not known

**How many parameters in the model below?**

- A) <12 B) 12-16 C) 17-20 D) more than 20
- How does this change if a variable can take on 3 values?

**What is conditionally independent of what?**



$P(A), P(B), P(C)$ : 3 parameters  
 $P(D|A,B,C)$ : 8  $P(E|C)$ : 2  $P(F|D,E)$ : 4  
 3+8+2+4=17

A trail  $\{X_1, X_2, \dots, X_k\}$  in the graph (no cycles) is an **active trail** if for each consecutive triplet in the trail:


- $X_{i-1} \rightarrow X_i \rightarrow X_{i+1}$ , and  $X_i$  is not observed
- $X_{i-1} \leftarrow X_i \leftarrow X_{i+1}$ , and  $X_i$  is not observed
- $X_{i-1} \rightarrow X_i \leftarrow X_{i+1}$ , and  $X_i$  is observed or one of its descendants is observed

**Variables connected by active trails are not conditionally independent**

- RL is hard**:
- never see result of action not taken
  - never told what best action was
  - often long sequence of actions before knowing their consequence

**types**:

- model-based**: learn full env ( $p(s'|s,a)$  &  $r(s,a)$ ), then DP or search (ex: chess, grid world with known transitions)
- model-free**: learn V-values, Q-values or policy (ex: Trading bots/online bidding, Atari from pixels)



**TD(0) update**:

discount factor  $\gamma$  one-step lookahead

$$V(s) \leftarrow V(s) + \alpha (r + \gamma V(s') - V(s))$$

learning rate  $\alpha$  immediate reward

**SARSA**: (on-policy)

$$Q(s,a) \leftarrow Q(s,a) + \alpha (r + \gamma Q(s',a') - Q(s,a))$$

uses next action taken

**Q-Learning**: (off-policy)

$$Q(s,a) \leftarrow Q(s,a) + \alpha (r + \gamma \max_{a'} Q(s',a') - Q(s,a))$$

uses optimal future action, learn the optimal policy even while exploring

**MDP** =  $(S, A, p(s'|s,a), r(s,a,s'), \gamma)$

value functions:  $V_\pi(s) = \mathbb{E}[G_t | S_t = s]$

$$Q_\pi(s,a) = \mathbb{E}[G_t | S_t = s, A_t = a]$$

$\Rightarrow V_\pi(s) = \sum \pi(a|s) Q_\pi(s,a)$

**Bellman's Eq**

- $V_\pi(s) = \sum_a \pi(a|s) \sum_{s'} p(s'|s,a) [r + \gamma V_\pi(s')]$
- $V^*(s) = \max_a \sum_{s'} p(s'|s,a) [r + \gamma V^*(s')]$
- $Q^*(s,a) = \sum_{s'} p(s'|s,a) [r + \gamma \max_{a'} Q^*(s',a')]$

**policy iteration**: (1) compute  $V_\pi$  using  $\pi$ ; (2)  $\pi'(s) = \arg \max_a Q_\pi(s,a)$


policy eval policy improvement

**exploration v/s exploitation strategy**:

**$\epsilon$ -greedy**:  $P(\text{explore}) = \epsilon = 1 - P(\text{exploit})$   $\hookrightarrow$  "on policy"

**MC RL**:  $V(s)$  = average of actual returns from episodes

## Overview of RL Strategies



## Take-aways (1)

- Superhuman game playing using DQN**
  - interestingly, superhuman play still requires search
  - current RL does not generalize well to new games
- Many combos of learning policy, models, V, Q-functions**
  - Separate policy selection from policy evaluation
- Sometimes use "behavioral cloning" to initialize**
- Often regularize**
  - To stabilize learning
  - To stay closer to previous policies or human play
- Increasingly use a more abstract state and action space**

## Take-aways (2)

Policy gradients work well for alignment

RLHF uses RL to propagate reward at end of the response (estimated by a model) back through the tokens in the response

Lots of tricks to support convergence: clipping, regularization, use of references

Today's methods: on policy, model free

Alternate: DPO directly optimizes the reward difference between a pair of responses

## What you should know


- Active learning**
- Uncertainty sampling  $\rightarrow$  Sample where models is uncertain
  - Query by committee  $\rightarrow$  Sample where model disagree
  - Information-based loss functions  $\rightarrow$  Maximize information gain  $\max KL \text{dis}(\cdot) \Rightarrow \max \mathbb{E}[I_O]$
- Optimal experimental design**
- A-optimal design  $\rightarrow$  Minimize trace  $= \sum \lambda_i$  of  $(X^T X)^{-1}$
  - D-optimal design  $\rightarrow$  Minimize det  $= \prod \lambda_i$  of  $(X^T X)^{-1}$
  - E-optimal design  $\rightarrow$  Minimize largest eigenvalue  $\max \lambda_i$  of  $(X^T X)^{-1}$
  - All minimize the variance in estimate of  $w$
- Response surface methods**  $\rightarrow$  Sequential experiments for optimization



$X = UD^T$ ; eigenvectors of  $X^T X$  = cols of  $V$ . eigvals of  $X^T X$  = squares of singular values of  $X$ .  $^2$  PCs are orthogonal  
 $^3$  PC scores of  $X$  are columns of  $Z = UD$ .  $^4$  If  $X$  is mean-centered, the PCA finds eigenvectors of  $X^T X$ , right singular vectors of  $X$  and projection directions of max covariance of  $X$ .  $^5$  Scores  $Z = XV = UD$  tell you where each sample lies along each PC; loading  $V$  tell you which feature contribute to each PC.  $^6$  Autoencoders generalize PCA  $^7$  or ICA!; and PCA = linear autoencoder with  $L_2$  loss  
 $^8$  PCA finds direction of max variance, ICA finds direction where data is most non-Gaussian/statistically indep.  $^9$  AEs learn weights to minimize reconstruction error.  $^{10}$  Real-world data often lies near a low-dimensional manifold, AEs (and PCA) try to learn coords on this manifold.  $^{11}$  Stacking AEs gives embeddings to use in SL.  $^{12}$  It is recommended to mean-center  $X$  before running PCA on it.  $^{13}$  The distributional similarity hypothesis for word embeddings states: Words occurring in similar contexts tend to have similar meanings.  $^{14}$  Relative to PCA, ICA aims to find PCs that are as statistically indep. as possible.  $^{15}$  the non-zero eigenvalues of  $AA^T$  and  $A^T A$  are the same.  $^{16}$  the left (right) singular vectors of a matrix  $A$  are the eigenvectors of  $AA^T$  ( $A^T A$ ).  $^{17}$  standardizing = mean-centering +  $\div$  by std  
 $^{18}$  mean encoding: replace each category with mean of target variable for that category. • Power methods for estimating eigenvectors are attractive because they guarantee to find a global optimum in reconstruction error when used in PCA  
 2018

$^{19}$  (MAYBE) best way to deal with missing data is to replace it by 0 and add an extra column indicating whether data is missing or not.  $^{20}$  Q-learning has higher bias and lower variance than MC methods.  $^{21}$  during PCA on a standardized dataset doesn't mean the scores will have variance = 1.  $^{22}$  Logistic reg gives a globally optimal solution to its loss function (binary cross entropy is convex); unlike NNets or k-means.  $^{23}$  In a penalized model, the less data we have, the more the overfitting risk, so the larger the regularization penalty selected by cross validation.  $^{24}$  degree  $d$  polynomial kernel SVM is more computationally expensive than linear kernel SVM, which is as expensive as logistic regression.  $^{25}$  A single perceptron update does not guarantee fixing this misclassification of the very point that triggered the update.  $^{26}$  In Adaboost for binary classification, training error is guaranteed to approach zero as the number of iteration  $\rightarrow \infty$ , and it should ideally use an underfit model as the "weak learner".  $^{27}$  For GMMs, complexity  $\uparrow$  as # Gaussians  $\uparrow$ , not the case for KNN with  $k \uparrow$ , nor NNets with  $L_2$  penalty coeff  $\uparrow$   
 2018

• hinge loss is convex! exponential loss too! 0-1 loss is non convex!  $^{28}$  Removal of a support vector doesn't always change the SVM decision boundary.  $^{29}$  In LR, looking at the largest regression weights in absolute value is NOT a reliable way to determine which features are most influential, UNLESS ALL FEATURES ARE STANDARDIZED!  $^{30}$  CNNs don't work well on high dimensional data like medical diagnosis from health records (with features like age, weight, temp, ...).  $^{31}$  the statement: "Vanilla RNNs, unlike HMMs, don't forget things exponentially quickly" is FALSE! • AEs don't ALWAYS reconstruct to a smaller dim  
ex: "Overcomplete AEs". • LMS does SGD in a NLL  $\uparrow$



scale inv: Linear regression (OLS), Decision trees,  $L_0$ , hard SVMs?  
non scale inv: Ridge (unless features are std),  $L_1$ , k-NN, PCA, SVMs  
coord-based: Gradient tree boosting  
not coord-based: NNs, Linear / Logistic regression

